

The process of obtaining a new image was performed by applying an affine transformation matrix randomly generated to the image given in this section.

Make sure you have set up OpenCV and NumPy for use in serial and parallel implementation, and CUDA for use in parallel implementation.

Sample images you can use are given in the dataset folder. The outputs obtained at the end of the program are saved in the output file.

You can use the following command to run the affineDeformation.py file:

python3 affineDeformation.py -f fileName.xxx

Note: fileName.xxx is the input image for which you want to apply the affine transformation process.

When you run it without any problems, you should get an output as follows:

```
(base) derya@derya:~/Desktop/affineDeformation/dataset$ python3 affineDeformation.py -f 3.pgm
->Loading images
    3.pgm successfully read!
Runtime of the program is 0.0056607723236083984
(base) derya@derya:~/Desktop/affineDeformation/dataset$
```

You can use the following command to run the affineDeformationParallel.py file:

sudo nvprof python3 affineDeformationParallel.py -f fileName.xxx

In this part, If you get an error regarding the installation of pycuda while running this command, you can run the following command for installation:

sudo pip install pycuda

When you run it without any problems, you should get an output as follows:

```
(base) derya@derya:~/Desktop/affineDeformation/dataset$ sudo nvprof python3 affineDeformationParallel.py -f 3.pgm
[sudo] password for derya:
==9744== NVPROF is profiling process 9744, command: python3 affineDeformationParallel.py -f 3.pgm
->Loading images
  3.pgm successfully read!
Runtime of the program is 0.0010502338409423828
==9744== Profiling application: python3 affineDeformationParallel.py -f 3.pgm
==9744== Profiling result:
   Type  Time(%)    Time       Calls      Avg       Min       Max    Name
GPU activities:  82.44%  182.79us        1  182.79us  182.79us  182.79us  affineDeformation
                  15.76%   34.944us        1   34.944us  34.944us  34.944us  [CUDA memcpy DtoH]
                  1.80%   4.0000us        2   2.0000us  1.4400us  2.5600us  [CUDA memcpy HtoD]
API calls:      77.10%  111.91ms        1  111.91ms  111.91ms  111.91ms  cuCtxCreate
                  22.31%   32.378ms        1   32.378ms  32.378ms  32.378ms  cuCtxDetach
                  0.17%   239.81us        1   239.81us  239.81us  239.81us  cuModuleLoadDataEx
                  0.13%   182.89us        1   182.89us  182.89us  182.89us  cuCtxSynchronize
                  0.12%   175.62us        3   58.541us  4.7630us  166.06us  cuMemAlloc
                  0.06%   82.120us        3   27.373us  3.0520us  73.769us  cuMemFree
                  0.05%   79.678us        1   79.678us  79.678us  79.678us  cuMemcpyDtoH
                  0.03%   41.410us        1   41.410us  41.410us  41.410us  cuModuleUnload
                  0.01%   20.618us        2   10.309us  5.8040us  14.814us  cuMemcpyHtoD
                  0.01%   18.654us        1   18.654us  18.654us  18.654us  cuLaunchKernel
                  0.01%   10.496us        1   10.496us  10.496us  10.496us  cuDeviceGetPCIBusId
                  0.00%   2.9570us        3      985ns    177ns    1.7800us  cuDeviceGetCount
                  0.00%   2.1640us        2   1.0820us    798ns    1.3660us  cuCtxGetDevice
                  0.00%   1.8110us        1   1.8110us  1.8110us  1.8110us  cuCtxPopCurrent
                  0.00%   1.3770us        2      688ns    642ns    735ns    cuDeviceGet
                  0.00%   1.3220us        3      440ns    225ns    713ns    cuDeviceGetAttribute
                  0.00%    784ns        1      784ns    784ns    784ns    cuModuleGetFunction
                  0.00%    761ns        1      761ns    761ns    761ns    cuDeviceComputeCapability
                  0.00%    652ns        1      652ns    652ns    652ns    cuCtxPushCurrent
                  0.00%    621ns        1      621ns    621ns    621ns    cuFuncSetBlockShape
(base) derya@derya:~/Desktop/affineDeformation/dataset$
```

You can also use --print-gpu-trace. You can use the following command by running:

sudo nvprof --print-gpu-trace python3 affineDeformationParallel.py -f fileName.xxx

When you run it without any problems, you should get an output as follows:

```
(base) derya@derya:~/Desktop/affineDeformation/dataset$ sudo nvprof --print-gpu-trace python3 affineDeformationParallel.py -f 3.pgm
==14742== NVPROF is profiling process 14742, command: python3 affineDeformationParallel.py -f 3.pgm
->Loading images
  3.pgm successfully read!
Runtime of the program is 2.0144617557525635
==14742== Profiling application: python3 affineDeformationParallel.py -f 3.pgm
==14742== Profiling result:
   Start  Duration      Grid Size    Block Size    Regs*    SSMem*    DSMem*    Size  Throughput  SrcMemType  DstMemType    De
vice  Context  Stream  Name
300.46ms  1.4080us        -          -          -          -          -    24B  16.256MB/s  Pageable    Device  GeForce MX1
10 (    1          7  [CUDA memcpy HtoD]
300.47ms  2.5920us        -          -          -          -          -  12.000KB  4.4152GB/s  Pageable    Device  GeForce MX1
10 (    1          7  [CUDA memcpy HtoD]
300.50ms  193.05us        (13 10 1)  (32 32 1)    28        0B        0B          -          -          -          -  GeForce MX1
10 (    1          7  affineDeformation [23]
300.71ms  37.087us        -          -          -          -          -  121.21KB  3.1169GB/s  Device      Pageable  GeForce MX1
10 (    1          7  [CUDA memcpy DtoH]

Regs: Number of registers used per CUDA thread. This number includes registers used internally by the CUDA driver and/or tools and can be more than what the compiler shows.
SSMem: Static shared memory allocated per CUDA block.
DSMem: Dynamic shared memory allocated per CUDA block.
SrcMemType: The type of source memory accessed by memory operation/copy
DstMemType: The type of destination memory accessed by memory operation/copy
(base) derya@derya:~/Desktop/affineDeformation/dataset$
```

If you do not want a detailed output like the one above, you can run the following command:

python3 affineDeformationParallel.py -f fileName.xxx